

# Generation of Human Faces

Author: Tommaso Ancilli

DIISM

# Preprocessing of Dataset

- Cleaning of the images:
  - I remove noise, being present on edges of the image, by cropping it. Going from a resolution of 216x180 to a 140x140
  - Down-sampling the image, going from 140x140 to a 64x64
- Feature scaling:
  - None
- Feature selection:
  - Randomly, each image was chosen to be inside the final dataset with a  $p$  probability

# Discriminator Architecture

```
1 -> Conv2d(3, 32, kernel_size=(7, 7), stride=(1, 1), padding=same)
2 -> AvgPool2d(kernel_size=2, stride=2, padding=0)
3 -> LeakyReLU(negative_slope=0.01, inplace=True)

4 -> Conv2d(32, 64, kernel_size=(5, 5), stride=(1, 1), padding=same)
5 -> AvgPool2d(kernel_size=2, stride=2, padding=0)
6 -> LeakyReLU(negative_slope=0.01, inplace=True)

7 -> Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
8 -> ReLU(inplace=True)

9 -> Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
10 -> ReLU(inplace=True)

11 -> Flatten(start_dim=1, end_dim=-1)
12 -> Linear(in_features=4096, out_features=1, bias=True)
13 -> Sigmoid()
```

Classification task: Sigmoid function for output layer

Downsampling the tensor

- Average Pooling
- Zero-padding convolution

# Discriminator loss function

Cost function:

- $-\frac{1}{n} * \sum_{i=1}^n (y_i \log(D(x)) + (1 - y_i) \log(1 - D(G(z))))$

Gradient-base optimizer: ADAM

- Learning rate :  $1 \times 10^{-5}$ ;  $\beta_1 = 0.9$ ;  $\beta_2 = 0.999$
- Minimization of the BCE

# Generator Architecture

Figure: "Paper-based" Generative architecture

```
1 -> ConvTranspose2d(5, 128, kernel_size=(4, 4), stride=(1, 1), padding=(2, 2))
2 -> ReLU()

3 -> BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
4 -> ConvTranspose2d(128, 64, kernel_size=(3, 3), stride=(3, 3), padding=(2, 2))
5 -> ReLU()

6 -> BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
7 -> ConvTranspose2d(64, 32, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))
8 -> ReLU()

9 -> BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
10 -> ConvTranspose2d(32, 16, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))
11 -> ReLU()

12 -> ConvTranspose2d(16, 3, kernel_size=(2, 2), stride=(2, 2))
13 -> ReLU()
```

Figure: Autoencoder for Generative architecture

```
1 -> Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
2 -> LeakyReLU(negative_slope=0.01, inplace=True)

3 -> Conv2d(64, 128, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
4 -> LeakyReLU(negative_slope=0.01, inplace=True)

5 -> Conv2d(128, 256, kernel_size=(7, 7), stride=(1, 1), padding=(3, 3))
6 -> LeakyReLU(negative_slope=0.01, inplace=True)

7 -> Conv2d(256, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
8 -> ReLU(inplace=True)

9 -> Conv2d(64, 3, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
10 -> ReLU(inplace=True)
```

## Problem faced

Auto-encoder has too many parameters and with the same learning rate for Discriminator and Generator the former was better to distinguish between real/fake images. Therefore, the min-max game ends up with the perfect separation of real from make-up images by the discriminator, preventing the weights update.

I have used two different solutions::

Changed architecture, I have found a "sweet-spot" between the ratio of the two learning rate following this paper (here). Concerning the objective function, reading the original paper (here)

# Solution for the Generator problem

1 attempt:

- $-\frac{1}{L} * \sum_{i=1}^L (y_i \log(D(x)) + (1 - y_i) \log(1 - D(G(z)))) =$   
 $-\frac{1}{L} * \sum_{i=1}^L (1 - y_i) \log(1 - D(G(z)))$
- Gradient-base optimizer: ADAM; Learning rate :  $1 \times 10^{-3}$ ;  $\beta_1 = 0.9$ ;  $\beta_2 = 0.999$
- Maximization of the objective

2 attempt:

- $-\frac{1}{L} * \sum_{i=1}^L (y_i \log(D(x)) + (1 - y_i) \log(1 - D(G(z)))) = -\frac{1}{L} * \sum_{i=1}^L y_i \log(D(G(z)))$
- Gradient-base optimizer: ADAM; Learning rate :  $1 \times 10^{-3}$ ;  $\beta_1 = 0.9$ ;  $\beta_2 = 0.999$
- Minimization of the objective

# Results for conventional architecture

Figure: Modified BCE





Figure: classic min-max problem with BCE





# Results for auto-encoder architecture

Figure: Modified BCE








Figure: classic min-max problem with BCE.(shortened to 4 epochs)



# Generated faces



# Conclusion

- Training process too unstable and the generation of the images is not straightforward
- Limited statistical variability in the image generation procedure